

LAB - Transform data using mapping data flows

In this tutorial, you'll use the Azure Data Factory user interface (UX) to create a pipeline that copies and transforms data from an Azure Data Lake Storage (ADLS) Gen2 source to an ADLS Gen2 sink using mapping data flow. The configuration pattern in this tutorial can be expanded upon when transforming data using mapping data flow

In this tutorial, you do the following steps:

- Create a data factory.
- Create a pipeline with a Data Flow activity.
- Build a mapping data flow with four transformations.
- Test run the pipeline.
- Monitor a Data Flow activity

Prerequisites

- **Azure subscription.** If you don't have an Azure subscription, create a [free Azure account](#) before you begin.
- **Azure storage account.** You use ADLS storage as a *source* and *sink* data stores. If you don't have a storage account, see [Create an Azure storage account](#) for steps to create one.

The file that we are transforming in this tutorial is MoviesDB.csv, which can be found [here](#). To retrieve the file from GitHub, copy the contents to a text editor of your choice to save locally as a .csv file. To upload the file to your storage account, see [Upload blobs with the Azure portal](#). The examples will be referencing a container named 'sample-data'.


Create a data factory

In this step, you create a data factory and open the Data Factory UX to create a pipeline in the data factory.

1. Open **Microsoft Edge** or **Google Chrome**. Currently, Data Factory UI is supported only in the Microsoft Edge and Google Chrome web browsers.
2. On the left menu, select **Create a resource > Integration > Data Factory**:

[Home](#) >

New

 Search the Marketplace

Azure Marketplace [See all](#)

Featured [See all](#)

Get started

Recently created

AI + Machine Learning

Analytics

Blockchain

Compute

Containers

Databases

Developer Tools

DevOps

Identity

Integration

Internet of Things

IT & Management Tools

Media

Migration

Mixed Reality

Monitoring & Diagnostics

Networking

Security

Software as a Service (SaaS)

Storage

Web



Logic App

[Quickstarts + tutorials](#)



API Management

[Quickstarts + tutorials](#)



Service Bus

[Quickstarts + tutorials](#)



Integration Account

[Quickstarts + tutorials](#)



Integration Service Environment

[Learn more](#)



Logic Apps Custom Connector

[Learn more](#)



Data Factory

[Quickstarts + tutorials](#)



Data Catalog

[Learn more](#)



Apache Kafka® on Confluent Cloud™
for Azure (preview)

[Learn more](#)



Dell Boomi Atom (Windows)
(preview)

[Learn more](#)

3. On the **New data factory** page, under **Name**, enter **ADFTutorialDataFactory**.

The name of the Azure data factory must be *globally unique*. If you receive an error message about the name value, enter a different name for the data factory. (for example, yournameADFTutorialDataFactory). For naming rules for Data Factory artifacts, see [Data Factory naming rules](#).

Create Data Factory ...

Basics Git configuration Networking Advanced Tags Review + create

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ

Resource group * ⓘ

[Create new](#)

Instance details

Region * ⓘ

Name * ⓘ

✖ The Data Factory name is already taken. Choose a different name.

Version * ⓘ

4. Select the Azure **subscription** in which you want to create the data factory.
5. For **Resource Group**, take one of the following steps:
 - a. Select **Use existing**, and select an existing resource group from the drop-down list.
 - b. Select **Create new**, and enter the name of a resource group.To learn about resource groups, see [Use resource groups to manage your Azure resources](#).
6. Under **Version**, select **V2**.
7. Under **Location**, select a location for the data factory. Only locations that are supported are displayed in the drop-down list. Data stores (for example, Azure Storage and SQL Database) and computes (for example, Azure HDInsight) used by the data factory can be in other regions.
8. Select **Create**.
9. After the creation is finished, you see the notice in Notifications center. Select **Go to resource** to navigate to the Data factory page.
10. Select **Author & Monitor** to launch the Data Factory UI in a separate tab.

Create a pipeline with a Data Flow activity

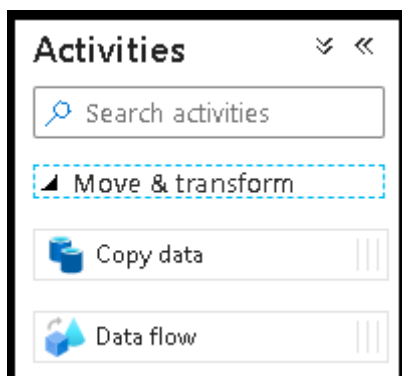
In this step, you'll create a pipeline that contains a Data Flow activity.

1. On the home page of Azure Data Factory, select **Orchestrate**.



2. In the **General** tab for the pipeline, enter **TransformMovies** for **Name** of the pipeline.

3. In the **Activities** pane, expand the **Move and Transform** accordion. Drag and drop the **Data Flow** activity from the pane to the pipeline canvas.



4. In the **Adding Data Flow** pop-up, select **Create new Data Flow** and then name your data flow **TransformMovies**. Click Finish when done.

A screenshot of the 'Adding Data Flow' dialog box. It has a title bar with a close button (X). Inside, there are two radio buttons: 'Use existing Data Flow' and 'Create new Data Flow'. The 'Create new Data Flow' option is selected. Below the radio buttons, there is a text input field labeled 'Data Flow name *' with a red asterisk indicating it is required. The text 'dataflow1' is entered into this field.

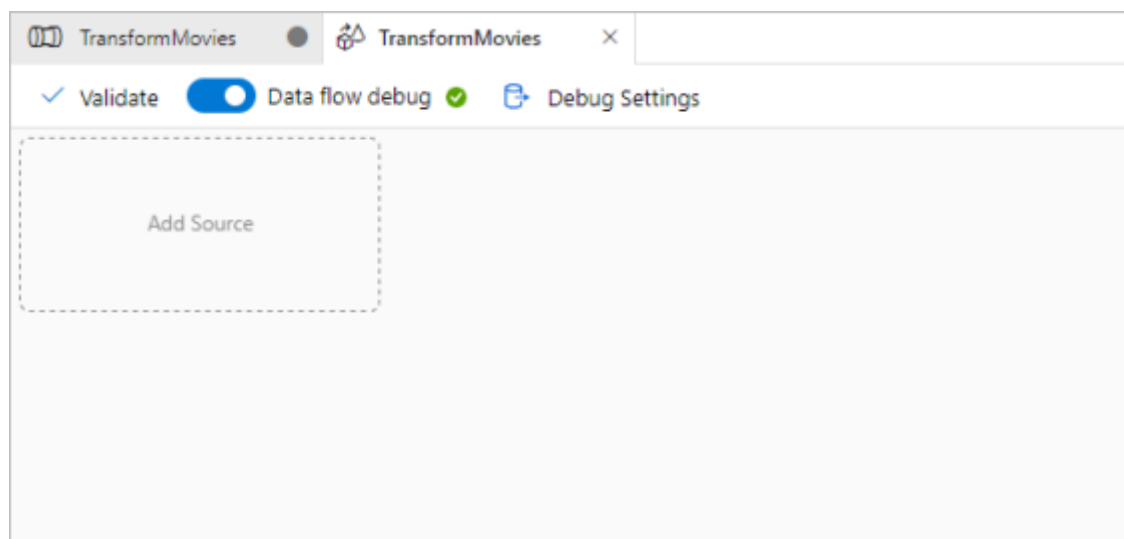
5. In the top bar of the pipeline canvas, slide the **Data Flow debug** slider on. Debug mode allows for interactive testing of transformation logic against a live Spark cluster. Data Flow clusters take 5-7 minutes to warm up and users are recommended to turn on debug first if they plan to do Data Flow development. For more information, see [Debug Mode](#).



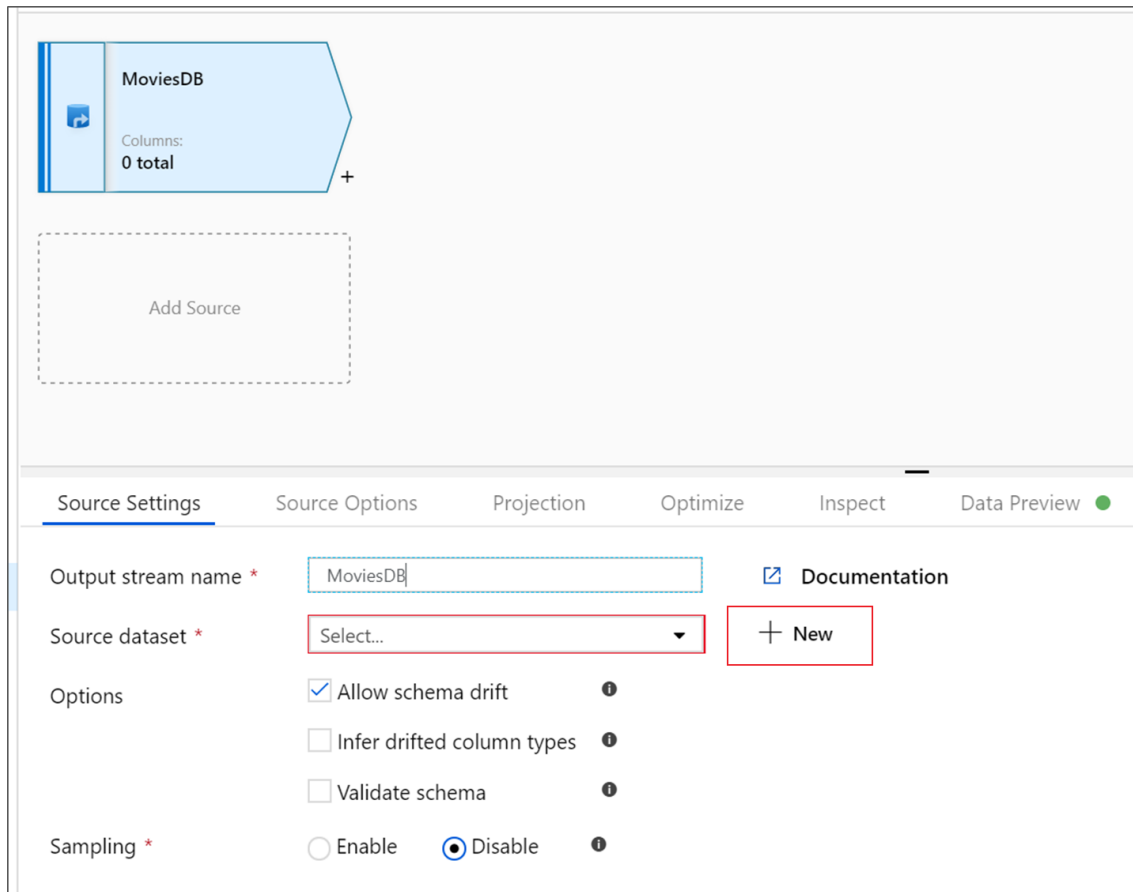
Build transformation logic in the data flow canvas

Once you create your Data Flow, you'll be automatically sent to the data flow canvas. In this step, you'll build a data flow that takes the moviesDB.csv in ADLS storage and aggregates the average rating of comedies from 1910 to 2000. You'll then write this file back to the ADLS storage.

1. In the data flow canvas, add a source by clicking on the **Add Source** box.



2. Name your source **MoviesDB**. Click on **New** to create a new source dataset.



MoviesDB
Columns:
0 total

Add Source

Source Settings | Source Options | Projection | Optimize | Inspect | Data Preview

Output stream name * [Documentation](#)

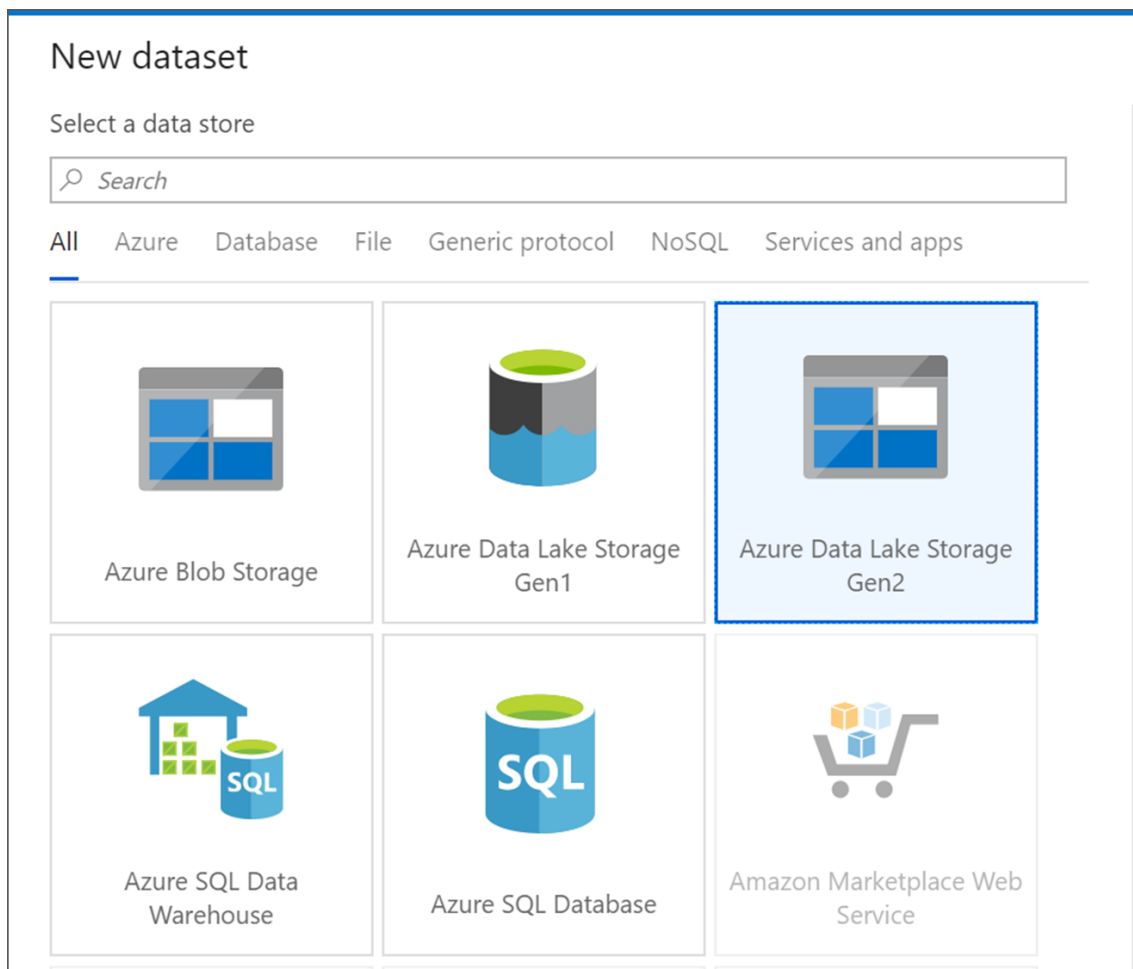
Source dataset * **+ New**

Options

- ☒ Allow schema drift ⓘ
- ☐ Infer drifted column types ⓘ
- ☐ Validate schema ⓘ

Sampling * ☐ Enable ☒ Disable ⓘ

3. Choose **Azure Data Lake Storage Gen2**. Click Continue.



New dataset

Select a data store







All | Azure | Database | File | Generic protocol | NoSQL | Services and apps

- Azure Blob Storage
- Azure Data Lake Storage Gen1
- Azure Data Lake Storage Gen2**
- Azure SQL Data Warehouse
- Azure SQL Database
- Amazon Marketplace Web Service

4. Choose **DelimitedText**. Click Continue.

Select format

Choose the format type of your data

 Parquet	 DelimitedText	 Json
 Avro	 ORC	 Binary

5. Name your dataset **MoviesDB**. In the linked service dropdown, choose **New**.

Set properties

Name
MoviesDB

Linked service *

Select...

Filter...

Select...

+ New

6. In the linked service creation screen, name your ADLS gen2 linked service **ADLSGen2** and specify your authentication method. Then enter your connection credentials. In this tutorial, we're using Account key to connect to our storage account. You can click **Test connection** to verify your credentials were entered correctly. Click Create when finished.

New linked service (Azure Data Lake Storage Gen2)

Name *

ADLSGen2

Description

Connect via integration runtime *

AutoResolveIntegrationRuntime

Authentication method

Account key

Account selection method

☒ From Azure subscription

☐ Enter manually

Azure subscription

Select all

Storage account name *

Test connection

☒ To linked service

☐ To file path

If the identity you use to access the data store only has permission to subdirectory instead of the entire account, specify the path to test connection. Please make sure your self-hosted integration runtime is higher than version 4.0 if connecting via self-hosted integration runtime.

Annotations

+ New

► Advanced

Create

 Test connection

Cancel

- Once you're back at the dataset creation screen, enter where your file is located under the **File path** field. In this tutorial, the file moviesDB.csv is located in container sample-data. As the file has headers, check **First row as header**. Select **From connection/store** to import the header schema directly from the file in storage. Click OK when done.

Set properties

Name

Linked service ^{*}

[Edit connection](#)

File path
 / / [Browse](#)

First row as header ☒

Import schema
☒ From connection/store ☐ From sample file ☐ None

► Advanced

8. If your debug cluster has started, go to the **Data Preview** tab of the source transformation and click **Refresh** to get a snapshot of the data. You can use data preview to verify your transformation is configured correctly.

Source Settings	Source Options	Projection	Optimize	Inspect	Data Preview ●	Description
Number of rows	+ INSERT 100	+ UPDATE 0	× DELETE 0	+ UPSERT 0	LOOKUP 0	TOTAL 9125
Typecast ▾	Modify ▾	Map drifted	Statistics	Remove		
movie	title	genres	year	Rating	Rotton Tomat	
108583	Fawty Towers (1975	Comedy	-1980	1	54	
32898	Trip to the Moon, A (Voyage ...	Action Adventure Fantasy Sci...	1902	7	80	
7243	Intolerance: Love's Struggle ...	Drama	1915	4	82	
62383	20,000 Leagues Under the Sea	Action Adventure Sci-Fi	1915	9	92	
8511	Immigrant, The	Comedy	1917	4	59	
3309	Dog's Life, A	Comedy	1917	3	83	

9. Next to your source node on the data flow canvas, click on the plus icon to add a new transformation. The first transformation you're adding is a **Filter**.

MoviesDB

Columns:
6 total

+

Row modifier

Filter

Add Source

10. Name your filter transformation **FilterYears**. Click on the expression box next to **Filter on** to open the expression builder. Here you'll specify your filtering condition.

Filter Settings Optimize Inspect Data Preview ●

Output stream name * [D](#)

Incoming stream * **MoviesDB**

Filter on *

Enter filter... ANY

11. The data flow expression builder lets you interactively build expressions to use in various transformations. Expressions can include built-in functions, columns from the input schema, and user-defined parameters. For more information on how to build expressions, see [Data Flow expression builder](#).

In this tutorial, you want to filter movies of genre comedy that came out between the years 1910 and 2000. As year is currently a string, you need to convert it to an integer using the `toInteger()` function. Use the greater than or equals to (`>=`) and less than or equals to (`<=`) operators to compare against literal year values 1910 and 2000. Union these expressions together with the and (`&&`) operator. The expression comes out as:

```
toInteger(year) >= 1910 && toInteger(year) <= 2000
```

To find which movies are comedies, you can use the `rlike()` function to find pattern 'Comedy' in the column genres. Union the rlike expression with the year comparison to get:

```
toInteger(year) >= 1910 && toInteger(year) <= 2000 && rlike(genres, 'Comedy')
```

If you've a debug cluster active, you can verify your logic by clicking **Refresh** to see expression output compared to the inputs used. There's more than one right answer on how you can accomplish this logic using the data flow expression language.

Visual Expression Builder [Expression reference documentation](#) ↗ ✕

FUNCTIONS «

All Functions Input schema Parameters

- abc movie
- abc title
- abc genres
- abc year
- abc Rating
- abc Rotten Tomato

`toInteger(year) >= 1910 && toInteger(year) <= 2000 && rlike(genres, 'Comedy')`

Data preview Refresh

Output: ↕	year abc	genres abc
✕	-1980	Comedy
✕	1902	Action Adventure Fantasy Sci-Fi
✕	1915	Drama War
✕	1915	Drama
✕	1915	Action Adventure Sci-Fi
✓	1917	Comedy
✓	1917	Comedy

[Save and finish](#) [Cancel](#) [Clear contents](#)

Click **Save and Finish** once you're done with your expression.

12. Fetch a **Data Preview** to verify the filter is working correctly.

Filter Settings

Optimize

Inspect

Data Preview

Description

Number of rows

+ INSERT

100

+ UPDATE

0

+ DELETE

0

+ UPSERT

0

LOOKUP

0

TOTAL 2021

Typecast

Modify

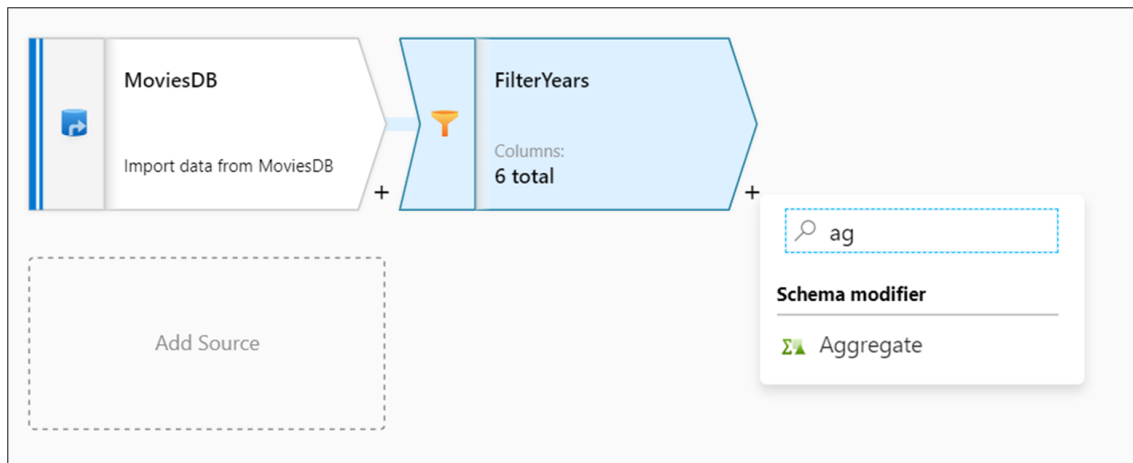
Map drifted

Statistics

Remove

movie	abc	title	abc	genres	abc	year	abc	Rating	abc	Rotten Tomat
+	8511	Immigrant, The		Comedy		1917		4		59
+	3309	Dog's Life, A		Comedy		1917		3		83
+	72626	Billy Blazes, Esq.		Comedy Western		1918		2		63
+	3310	Kid, The		Comedy Drama		1921		8		57
+	83096	Haunted House, The		Comedy		1920		9		61
+	83318	Goat, The		Comedy		1921		8		86
+	83322	Boat, The		Comedy		1920		4		65

13. The next transformation you'll add is an **Aggregate** transformation under **Schema modifier**.



14. Name your aggregate transformation **AggregateComedyRatings**. In the **Group by** tab, select **year** from the dropdown to group the aggregations by the year the movie came out.

Aggregate Settings		Optimize	Inspect	Data Preview
Output stream name *	AggregateComedyRating	Documentation		
Incoming stream *	FilterYears			
Group by		Aggregates		
FilterYears's column		Name as		
abc year		year		+ -

15. Go to the **Aggregates** tab. In the left text box, name the aggregate column **AverageComedyRating**. Click on the right expression box to enter the aggregate expression via the expression builder.

Aggregate Settings		Optimize	Inspect	Data Preview
Output stream name *	AggregateComedyRating	Documentation		
Incoming stream *	FilterYears			
Group by		Aggregates		
Grouped by: year				
AverageComedyRating		Enter expression...		ANY + -

16. To get the average of column **Rating**, use the `avg()` aggregate function. As **Rating** is a string and `avg()` takes in a numerical input, we must convert the value to a number via the `toInteger()` function. This expression looks like:

```
avg(toInteger(Rating))
```

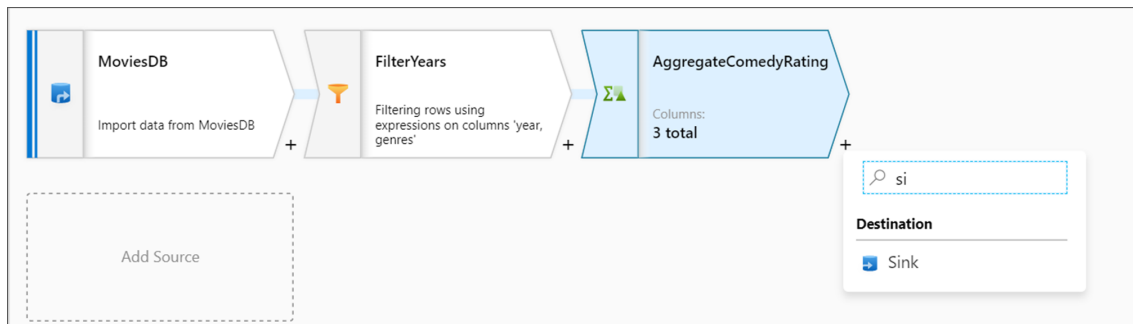
Click **Save and Finish** when done.

OUTPUT SCHEMA	FUNCTIONS	EXPRESSION FOR FIELD "AVERAGECOMEDYRATING"
ANY AverageComedyRating	<div>Filter...</div> <div>All Functions Input schema Parameters</div> <div>abc movie</div> <div>abc title</div>	<code>avg(toInteger(Rating))</code>

17. Go to the **Data Preview** tab to view the transformation output. Notice only two columns are there, **year** and **AverageComedyRating**.

Aggregate Settings	Optimize	Inspect	Data Preview
Output stream name * <input type="text" value="AggregateComedyRating"/> Documentation			
Incoming stream * FilterYears			
Group by Aggregates			
Grouped by: year			
<div><div>AverageComedyRating</div><div>Enter expression...</div><div>ANY</div><div>+</div><div>⌵</div></div>			

18. Next, you want to add a **Sink** transformation under **Destination**.



19. Name your sink **Sink**. Click **New** to create your sink dataset.







Sink	Settings	Mapping	Optimize	Inspect	Data Preview
Output stream name * <input type="text" value="Sink"/> Documentation					
Incoming stream * AggregateComedyRating					
Sink dataset * <div>Select...</div> <div>+ New</div>					
Options					
<input checked="" type="checkbox"/> Allow schema drift ⓘ					
<input type="checkbox"/> Validate schema ⓘ					

20. Choose **Azure Data Lake Storage Gen2**. Click Continue.

New dataset

Select a data store







[All](#)
[Azure](#)
[Database](#)
[File](#)
[Generic protocol](#)
[NoSQL](#)
[Services and apps](#)

 <p>Azure Blob Storage</p>	 <p>Azure Data Lake Storage Gen1</p>	 <p>Azure Data Lake Storage Gen2</p>
 <p>Azure SQL Data Warehouse</p>	 <p>Azure SQL Database</p>	 <p>Amazon Marketplace Web Service</p>

21. Choose **DelimitedText**. Click Continue.

Select format

Choose the format type of your data

 <p>Parquet</p>	 <p>DelimitedText</p>	 <p>Json</p>
 <p>Avro</p>	 <p>ORC</p>	 <p>Binary</p>

22. Name your sink dataset **MoviesSink**. For linked service, choose the ADLS gen2 linked service you created in step 6. Enter an output folder to write your data to. In this tutorial, we're writing to folder 'output' in container 'sample-data'. The folder doesn't need to exist beforehand and can be dynamically created. Set **First row as header** as true and select **None** for **Import schema**. Click Finish.

Set properties

Name
MoviesSink

Linked service *
ADLSGen2

[Edit connection](#)

File path
sample-data / output / File [Browse](#)

First row as header ☒

Import schema
☐ From connection/store
 ☐ From sample file
 ☒ None

▲ Advanced

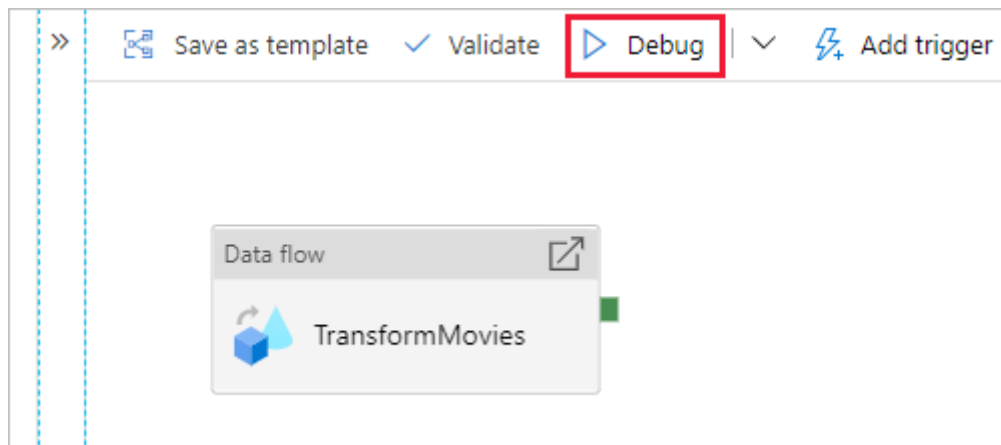
[Open this dataset](#) for more advanced configuration with parameterization.

Now you've finished building your data flow. You're ready to run it in your pipeline.

Running and monitoring the Data Flow

You can debug a pipeline before you publish it. In this step, you're going to trigger a debug run of the data flow pipeline. While data preview doesn't write data, a debug run will write data to your sink destination.

1. Go to the pipeline canvas. Click **Debug** to trigger a debug run.






2. Pipeline debug of Data Flow activities uses the active debug cluster but still take at least a minute to initialize. You can track the progress via the **Output** tab. Once the run is successful, click on the eyeglasses icon to open the monitoring pane.





General

Parameters

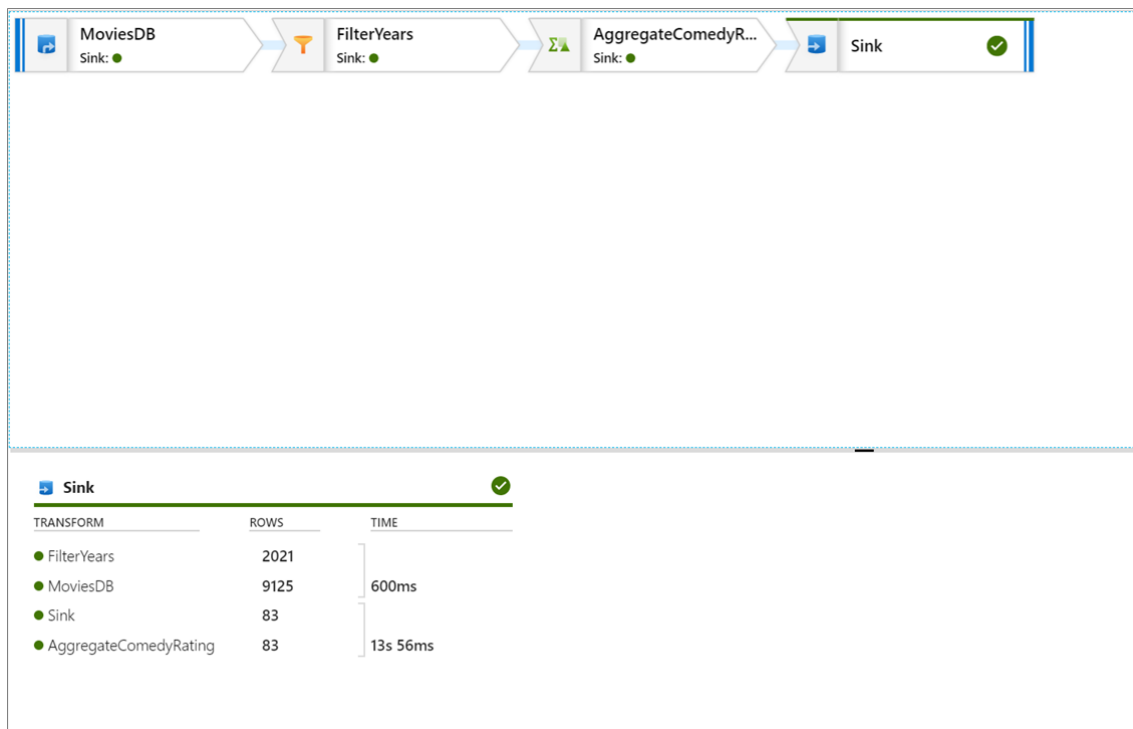
Variables

Output

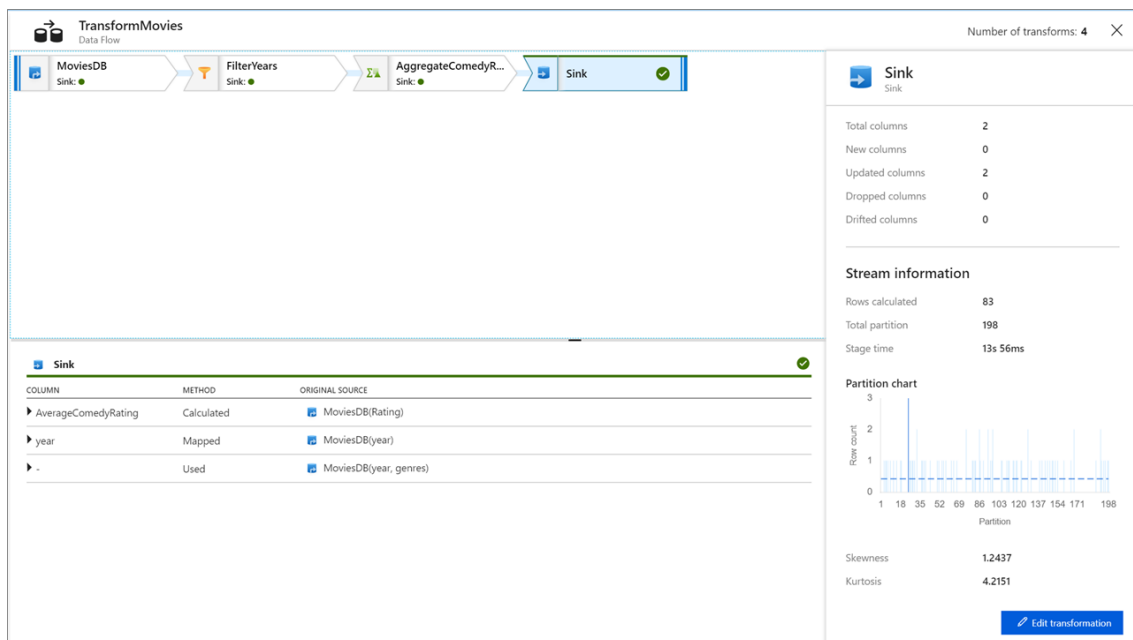
Pipeline run ID: **d67bdd03-4807-468b-abf2-09e9418660bf**   

NAME	TYPE	RUN START	DURATION	STATUS	ACTIONS	RUN ID
TransformMo...	ExecuteDataF...	10/03/2019 5:34 PM	00:02:03	 Succeeded	  	d7c9a07f-cb71-4e48-a016-e895c9ea96d5

3. In the monitoring pane, you can see the number of rows and time spent in each transformation step.



4. Click on a transformation to get detailed information about the columns and partitioning of the data.



If you followed this tutorial correctly, you should have written 83 rows and 2 columns into your sink folder. You can verify the data is correct by checking your blob storage.

Next steps

The pipeline in this tutorial runs a data flow that aggregates the average rating of comedies from 1910 to 2000 and writes the data to ADLS. You learned how to:

- Create a data factory.
- Create a pipeline with a Data Flow activity.
- Build a mapping data flow with four transformations.
- Test run the pipeline.
- Monitor a Data Flow activity

